

# Lecture 6: Working with data frames (and tibbles)

---

Ingmar Sturm

UCSB

2024-07-02

Special thanks to Robin Liu for select course content used with permission.

# Load the libraries

```
library(tidyverse)
library(gapminder)
```

Recall what's in `gapminder`

```
head(gapminder, 6)
```

```
## # A tibble: 6 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
```

# Filtering

Retrieve the data for Afghanistan for years after 1979.

```
## # A tibble: 6 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1982   39.9 12881816   978.
## 2 Afghanistan Asia      1987   40.8 13867957   852.
## 3 Afghanistan Asia      1992   41.7 16317921   649.
## 4 Afghanistan Asia      1997   41.8 22227415   635.
## 5 Afghanistan Asia      2002   42.1 25268405   727.
## 6 Afghanistan Asia      2007   43.8 31889923   975.
```

03:00

# Filtering

A much better way using `dplyr::filter()` (which is part of `tidyverse`).

```
filter(gapminder, country == "Afghanistan", year > 1979)
```

```
## # A tibble: 6 x 6
```

```
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1982   39.9 12881816    978.
## 2 Afghanistan Asia      1987   40.8 13867957    852.
## 3 Afghanistan Asia      1992   41.7 16317921    649.
## 4 Afghanistan Asia      1997   41.8 22227415    635.
## 5 Afghanistan Asia      2002   42.1 25268405    727.
## 6 Afghanistan Asia      2007   43.8 31889923    975.
```

Less repetition, easier to read.

# Selecting

```
head(gapminder)
```

```
## # A tibble: 6 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
```

Select only `country`, `year`, `lifeExp`, and `pop` variables.

02:00

# Selecting

## Easier way with `dplyr::select()`

```
head(select(gapminder, country, year, lifeExp, pop), 3)
```

```
## # A tibble: 3 x 4
##   country      year lifeExp      pop
##   <fct>      <int>   <dbl>   <int>
## 1 Afghanistan  1952    28.8  8425333
## 2 Afghanistan  1957    30.3  9240934
## 3 Afghanistan  1962    32.0 10267083
```

Also works with negative indices:

```
head(select(gapminder, -country, -year, -lifeExp, -pop), 3)
```

```
## # A tibble: 3 x 2
##   continent gdpPercap
##   <fct>      <dbl>
## 1 Asia      779.
## 2 Asia      821.
## 3 Asia      853.
```

# Combining select and filter

```
head(gapminder)
```

```
## # A tibble: 6 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>   <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952    28.8  8425333    779.
## 2 Afghanistan Asia      1957    30.3  9240934    821.
## 3 Afghanistan Asia      1962    32.0 10267083    853.
## 4 Afghanistan Asia      1967    34.0 11537966    836.
## 5 Afghanistan Asia      1972    36.1 13079460    740.
## 6 Afghanistan Asia      1977    38.4 14880372    786.
```

Return the `country`, `year`, and `lifeExp` for Rwanda in years between 1960 and 1970.

```
## # A tibble: 2 x 3
##   country  year lifeExp
##   <fct>   <int>   <dbl>
## 1 Rwanda  1962     43
## 2 Rwanda  1967    44.1
```

03:00

# The Pipe

The operator `|>` pipes the left-hand side as **the first** argument to the right-hand side, and returns the result.

```
25 |> sqrt() # no argument passed in explicitly
```

```
## [1] 5
```

```
increment_power <- function(x, pwr = 2) {  
  x <- x + 1  
  return(x^pwr)  
}
```

```
1 |> increment_power() # equiv. increment_power(1)
```

```
## [1] 4
```

```
1 |> increment_power(3) # equiv. increment_power(1, 3)
```

```
## [1] 8
```



# The Pipe

Use the pipe `|>` to simplify common operations.

```
select(filter(gapminder, country == "Rwanda", year > 1960, year < 1970),  
        country, year, lifeExp)
```

```
## # A tibble: 2 x 3  
##   country  year lifeExp  
##   <fct>   <int>   <dbl>  
## 1 Rwanda   1962     43  
## 2 Rwanda   1967    44.1
```

```
gapminder |>  
  filter(country == "Rwanda", year > 1960, year < 1970) |>  
  select(country, year, lifeExp)
```

```
## # A tibble: 2 x 3  
##   country  year lifeExp  
##   <fct>   <int>   <dbl>  
## 1 Rwanda   1962     43  
## 2 Rwanda   1967    44.1
```

# The Pipe

Another approach you may have tried:

```
filtered_gap <- filter(gapminder, country == "Rwanda", year > 1960, year < 1970)
select(filtered_gap, country, year, lifeExp)
```

```
## # A tibble: 2 x 3
##   country  year lifeExp
##   <fct>   <int>   <dbl>
## 1 Rwanda  1962     43
## 2 Rwanda  1967    44.1
```


```
gapminder |>
  filter(country == "Rwanda", year > 1960, year < 1970) |>
  select(country, year, lifeExp)
```

```
## # A tibble: 2 x 3
##   country  year lifeExp
##   <fct>   <int>   <dbl>
## 1 Rwanda  1962     43
## 2 Rwanda  1967    44.1
```

# The Pipe

```
gapminder |>  
  filter(country == "Rwanda", year > 1960, year < 1970) |>  
  select(country, year, lifeExp)
```

```
## # A tibble: 2 x 3  
##   country year lifeExp  
##   <fct>   <int>   <dbl>  
## 1 Rwanda  1962     43  
## 2 Rwanda  1967    44.1
```



## Piping in **R** is like baking

```
slice(decorate(bake(mix(ingredients))))
```



|>

mix( ) |>

bake( ) |>

decorate( ) |>

slice( ) ->

@ArthurWelle

# The Pipe

Return the *first 3 rows* of `continent`, `year`, and `lifeExp` of the "Americas" continent by rewriting the following code using pipes.

```
head(select(filter(gapminder, country == "Chile"), continent, year, lifeExp), 3)
```

```
## # A tibble: 3 x 3
##   continent year lifeExp
##   <fct>      <int>   <dbl>
## 1 Americas  1952    54.7
## 2 Americas  1957    56.1
## 3 Americas  1962    57.9
```

03:00

# dplyr and tidyverse

Reminder that `select()` and `filter()` are functions in the `dplyr` package which is part of the `tidyverse`. You must call `library(dplyr)` or `library(tidyverse)` to use it. `library(tidyverse)` automatically loads `dplyr`; it is sometimes called a *meta-package*.

Interestingly `|>` is a base R pipe.

`tidyverse` has its own pipe which looks like `%>%`; you might see this in the wild. It became so popular that `|>` was recently added to base R.

Which one should you use? It's a matter of preference but, in short, use `%>%` if you want greater flexibility and you want your code to work without the `tidyverse`. Use `|>` if you want speed. [comparison](#)

```
25 |> sqrt()
```

```
## [1] 5
```

```
25 %>% sqrt()
```

```
## [1] 5
```

# Comparing dplyr and base R

```
gapminder |>  
  filter(country == "Chile") |>  
  select(continent, year, lifeExp) |>  
  head(3)
```

```
## # A tibble: 3 x 3  
##   continent year lifeExp  
##   <fct>      <int>   <dbl>  
## 1 Americas  1952    54.7  
## 2 Americas  1957    56.1  
## 3 Americas  1962    57.9
```

```
head(gapminder[gapminder$country == "Chile", c("continent", "year", "lifeExp")], 3)
```

```
## # A tibble: 3 x 3  
##   continent year lifeExp  
##   <fct>      <int>   <dbl>  
## 1 Americas  1952    54.7  
## 2 Americas  1957    56.1  
## 3 Americas  1962    57.9
```

# More practice

Return a tibble containing `country`, `year`, and `gdpPercap` for countries with GDP per cap less than 300 for years before 2007.

```
## # A tibble: 3 x 3
##   country          year gdpPercap
##   <fct>          <int>     <dbl>
## 1 Congo, Dem. Rep.  2002      241.
## 2 Guinea-Bissau    1952      300.
## 3 Lesotho          1952      299.
```

05:00



# Creating a tibble (or a data frame)

```
(dogs <- tibble(name = c("Ralph", "Sully", "Capsule"),  
                age = c(3, 12, 1)))
```

```
## # A tibble: 3 x 2
```

```
##   name      age
```

```
##   <chr>   <dbl>
```

```
## 1 Ralph      3
```

```
## 2 Sully     12
```

```
## 3 Capsule    1
```

```
is_rescue <- c(F, F, T)
```

```
(dogs <- tibble(name = c("Ralph", "Sully", "Capsule"),  
                age = c(3, 12, 1),  
                rescue = is_rescue))
```

```
## # A tibble: 3 x 3
```

```
##   name      age rescue
```

```
##   <chr>   <dbl> <lgl>
```

```
## 1 Ralph      3 FALSE
```

```
## 2 Sully     12 FALSE
```

```
## 3 Capsule    1  TRUE
```

# Adding variables to an existing tibble

```
(dogs <- tibble(name = c("Ralph", "Sully", "Capsule"),  
                age = c(3, 12, 1)))
```

```
## # A tibble: 3 x 2
```

```
##   name      age
```

```
##   <chr>   <dbl>
```

```
## 1 Ralph      3
```

```
## 2 Sully     12
```

```
## 3 Capsule    1
```

```
is_rescue <- c(F, F, T)
```

```
dogs$rescue <- is_rescue
```

```
dogs
```

```
## # A tibble: 3 x 3
```

```
##   name      age rescue
```

```
##   <chr>   <dbl> <lgl>
```

```
## 1 Ralph      3 FALSE
```

```
## 2 Sully     12 FALSE
```

```
## 3 Capsule    1  TRUE
```

# Adding variables to an existing tibble

```
head(gapminder, 1)
```

```
## # A tibble: 1 x 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>   <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8 8425333    779.
```

Return the *total* GDP for each row. This is  $\text{pop} \times \text{gdpPercap}$ .

```
my_gap <- gapminder # create a copy so we don't overwrite gapminder
my_gap$totalGdp <- my_gap$pop * my_gap$gdpPercap
head(my_gap, 2)
```

```
## # A tibble: 2 x 7
##   country      continent  year lifeExp      pop gdpPercap  totalGdp
##   <fct>        <fct>    <int>  <dbl>   <int>    <dbl>    <dbl>
## 1 Afghanistan Asia      1952   28.8 8425333    779. 6567086330.
## 2 Afghanistan Asia      1957   30.3 9240934    821. 7585448670.
```

Remember `my_gap$pop` and `my_gap$gdpPercap` are vectors.

# dplyr::Mutate()

```
(dogs <- tibble(name = c("Ralph", "Sully", "Capsule"),  
  age = c(3, 12, 1)))
```

```
## # A tibble: 3 x 2  
##   name      age  
##   <chr>   <dbl>  
## 1 Ralph      3  
## 2 Sully     12  
## 3 Capsule    1
```

```
dogs <- tibble(name = c("Ralph", "Sully", "Capsule"),  
  age = c(3, 12, 1))
```

```
dogs |> mutate(rescue = c(F, F, T))
```

```
## # A tibble: 3 x 3  
##   name      age rescue  
##   <chr>   <dbl> <lgl>  
## 1 Ralph      3 FALSE  
## 2 Sully     12 FALSE  
## 3 Capsule    1  TRUE
```

# Adding Variables

Return a tibble containing `country`, `year`, and *total GDP* for countries with GDP per cap less than 300 for years before 2007.

```
gapminder |>
  mutate(totalGdp = pop * gdpPercap) |>
  filter(gdpPercap < 300, year < 2007) |>
  select(country, year, totalGdp)
```

```
## # A tibble: 3 x 3
##   country      year  totalGdp
##   <fct>      <int>    <dbl>
## 1 Congo, Dem. Rep. 2002 13355730548.
## 2 Guinea-Bissau   1952  174108987.
## 3 Lesotho        1952  223760205.
```

```
head(gapminder, 1) # Remember the original tibble is unmodified!!
```

```
## # A tibble: 1 x 6
##   country    continent  year lifeExp    pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>  <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8 8425333    779.
```

# Saving the result to a new tibble

```
my_gap <- gapminder |>
  mutate(totalGdp = pop * gdpPercap) |>
  filter(gdpPercap < 300, year < 2007) |>
  select(country, year, totalGdp)
```

```
my_gap
```

```
## # A tibble: 3 x 3
```

```
##   country      year  totalGdp
##   <fct>      <int>    <dbl>
## 1 Congo, Dem. Rep. 2002 13355730548.
## 2 Guinea-Bissau   1952  174108987.
## 3 Lesotho         1952  223760205.
```

# Summary

- It is common to manipulate tibbles and data frames
- `filter`, `select`, and `mutate` are useful functions in `dplyr` which is part of the `tidyverse` meta-package
- the pipe `|>` operator simplifies a lot of operations, but don't go overboard!!

$$\frac{5^7 - 2\sqrt{4}}{\log_2(100)}$$

```
(5^7 - 2*sqrt(4)) / log(100, base = 2)
```

```
## [1] 11758.38
```

```
5 %>% # I'm using %>% here because this is not allowed with |>
```

```
  '^'(7) %>%
```

```
  '-'(2 %>% '*'(4 %>% sqrt())) %>%
```

```
  '/'(100 %>% log(2))
```

```
## [1] 11758.38
```

# More on dplyr

`dplyr` contains many more functions for *data wrangling* (aka *data munging*).

See the [dplyr cheatsheet](#)

Most important things I left out: `group_by` and `summarize`.

```
gapminder |>
  group_by(continent) |>
  summarize(avg_life_exp = mean(lifeExp), avg_gdpPercap = mean(gdpPercap))
```

```
## # A tibble: 5 x 3
##   continent avg_life_exp avg_gdpPercap
##   <fct>      <dbl>      <dbl>
## 1 Africa      48.9      2194.
## 2 Americas    64.7      7136.
## 3 Asia        60.1      7902.
## 4 Europe      71.9     14469.
## 5 Oceania     74.3     18622.
```